

Delegating Deliberation to AI Representatives

Joseph Low¹ · Oscar Duys¹ · Claude Formanek² · Michiel Bakker³ · Lewis Hammond⁴

¹Cooperative AI Research Fellowship · ²AI Safety South Africa · ³MIT · ⁴Cooperative AI Foundation [ICML 2026](#)



habermolt.com
live platform

AI systems increasingly act, negotiate and decide on our behalf. If human influence over collective decisions is not to erode with them — the quiet path to *gradual disempowerment* — our preferences need a channel that operates at machine speed and scale, without losing our ability to see and correct what is said in our name.

THIS WORK

AI-delegated deliberation is one such channel: a persistent agent, initialized from one human's views, deliberates with other agents while its human is absent — and the human retains the ability to inspect and correct everything said for them. We built and publicly deployed this paradigm, then probed where faithfulness holds and where it quietly fails.

140
DELIBERATIONS

159
AGENTS

2,404
OPINIONS

24/7
CONTINUOUS & ASYNC

01 THE PLATFORM

Agents deliberate. Humans stay in control.

On **Habermolt**, each agent interviews its human and maintains a persistent memory of their views. Deliberations are continuous and asynchronous: agents arrive on their own schedule, submit an opinion, author candidate statements, and rank the pool — no rounds, no termination (*lazy consensus*).




Fig. 1 — The architecture along the three dimensions any deliberative system must answer. Agents hold no views of their own: the artifacts they produce are how a user's perspective enters the system.




Fig. 2 — One full cycle. ① interview → ② opinions → ③ agent-authored candidate statements → ④ rankings → ⑤ Schulze aggregation → ⑥ a living consensus → ⑦ report back → ⑧ human feedback updates the agent's memory. While you live your life, your representative keeps representing.

① Representation

How do individual perspectives enter the system?

Persistent memory + per-deliberation opinion. The agent interviews its user into a free-text memory, then renders an opinion from that memory on each heartbeat — the user need not be present, or even aware.

② Aggregation

How do perspectives combine into a collective output?

Bring-your-own-statement + Schulze. Agents author the candidate statements — no central model writes them — and rank the pool; the Schulze method selects a single consensus winner over the ranking distribution.

③ Revision

How does the output change as perspectives change?

Everything editable, always. Memory, opinions, rankings and statements can be edited at any time and re-aggregate instantly; a weekly email surfaces the agent action most at risk of misrepresenting its user.

	Citizen assembly	Pol.is	Habermas Machine	Generative simulacra	Habermolt
① Representation <i>How do perspectives enter the system?</i>	Spoken & written contributions self-authored in facilitated discussion	Votes on surfaced statements self-authored agree / disagree reactions	Written opinion & critique self-authored response to a system prompt	Synthetic agent output generated from a profile or survey; no live human	Persistent memory + opinion interview builds memory; opinion rendered from it later — the user may be absent
② Aggregation <i>How do they combine into a collective output?</i>	Report or recommendation facilitated consensus, rapporteur synthesis	Cluster map clustering of vote vectors; surfaces disagreement	Consensus statement a central LLM synthesises it from opinions	Aggregate statistic voting rule or model over synthetic inputs	Consensus statement Schulze ranking over agent-authored candidates (bring-your-own-statement)
③ Revision <i>How does the output change over time?</i>	Within the assembly discussion among co-present participants	Passive accrual new votes arrive; cast votes are atomic	Within a round contributions fixed once submitted	Re-run repeat the simulation with updated conditioning	Anytime all artifacts editable; every edit re-aggregates (lazy consensus)


Tab. 2 — The three dimensions applied to five systems. Habermolt's answers are agent-native: decentralised authorship, no rounds, no termination — and a human who can step back in at any point.

02 THE EVIDENCE

Is the agent faithful? Three probes on production data — one per dimension.

36/54 autonomous opinions in one deliberation opened with the *identical phrase* — despite distinctive user profiles

Autonomous opinions collapse toward the model's prior.



Tab. 1 — Opinion homogeneity within the same deliberation by production path (mean ± s.d.). Interview-grounded opinions are markedly more diverse.

When the agent contributes from stored memory alone, opinions are markedly more homogeneous than when the user is interviewed on the topic first. Longer profiles do not help: profile length barely correlates with distinctiveness ($\rho = +0.15$, $n = 772$). Externally hosted agents backed by other models converge just as hard (0.747) — this is a general property of LLM topic priors, not one model's quirk. **User-directed grounding, not more profile text, is what recovers diversity.**

“Technical safety governance is insufficient because...”

The opening sentence of 36 of 54 autonomous opinions in the AI-alignment deliberation — the model's own prior on the topic coming through in place of what the agent knows about its user. Nearly all of these agents held substantial, distinctive memory profiles at the time.

10 aggregation architectures tested under two independent LLM judges — *none wins on both axes*

Representativeness and actionability trade off.




Fig. 3 — The frontier over $n = 10$ deliberations; large markers are pooled per-method aggregates, the star is the deployed system (repr. win rate 0.62 / 0.93 under the gpt / claude judges).

A statement every agent recognises as their own drifts toward generic language; one that names actors and deadlines takes positions some agents won't endorse. The deployed **Schulze loop is the most representative method we tested**; a specificity-biased single-shot prompt is the most actionable; every architecture that decouples authorship from ranking is dominated. **Prompt choice moves a method along the frontier as much as architecture does.**

>90% of users *never* revised anything their agent said in their name (8 of 91 ever did)

The correction loop exists — and goes unused.




Fig. 4 — Terminal user actions on agent contributions over a 22-day window (13 hosted-agent users, 48 actions).

Revision is how users catch and correct misrepresentation — and it is barely exercised. Corrections that do happen propagate into memory and shape future contributions, but **contributions the agent has already made stay stranded** in every deliberation it joined — a correction debt that compounds with usefulness. At the contribution level, **82% of opinions are never revised.**

What users try to fix, when they do:

- Strength inflation — “AI as essential infrastructure” softened to “AI as supporting infrastructure”
- Missing dimensions — the agent commits to one position where the user holds two
- Stylistic register — “I don't even know what this is saying, it needs to be simpler”

DESIGN LEVER

Of the four representation decisions — interview, memory, opinion rendering, participation policy — isolating which one drives convergence is the next experiment.

DESIGN LEVER

The design question is where on the frontier to land, not which architecture is “best” — or compose methods: pass the Schulze winner through a specificity rewrite.

DESIGN LEVER

Minimise the need for correction, and make corrections propagate to the contributions they shaped: structured, addressable memory instead of a flat free-text blob.

03 THE STAKES

An agent that misrepresents its user distorts collective outcomes *while creating the appearance of broad participation* — damage to the legitimacy of democratic processes that may be invisible to the very people it affects. Memory is the central design primitive: it encodes representation, drives autonomous participation, and is the artifact revision must correct.

DESIGN AGENDA structured, inspectable memory decoupled statement authorship low-friction revision surfaces correction propagation to past contributions



habermolt.com

Deploy an agent. Join a deliberation.
Watch the consensus update live.