


Delegating Deliberation to AI Agents

An experience report from habermolt.com — what works, what doesn't, and where the leverage is.

Joseph Low · Oscar Duys · Cooperative AI Research Fellowship · 2026




Cooperative AI
Research Fellowship


THE QUESTION

Can AI agents represent humans well enough to deliberate on their behalf?

Democratic deliberation doesn't scale — classical methods cost hundreds of paid participant-hours and leave the people who show up not matching the people who have something to say. **Habermolt** is our deployed answer: humans teach AI agents their values through profiles and chat interviews, then deploy them into live, continuous deliberations that run 24/7. We analysed **71 production deliberations** with 1,000+ consensus statements to ask, candidly, what works.



THE ARCHITECTURE




THE BAD NEWS


Diversity compresses at every stage of the pipeline.

We probed three stages and found the same story at each. **Opinions** compress the profiles they are generated from. **Statements** compress the opinions they are supposed to summarise. **Rankings** carry enough noise to swamp small differences between near-identical candidates. Diversity is lost at every hand-off — and it does not come back downstream.

Diverse human views → homogeneous opinions → near-identical statements → indistinguishable actions.



Opinions compress profiles. Autonomous opinions (profile only) are significantly more homogeneous than interview-sourced opinions ($p < 10^{-6}$). 36 of 54 agents in the AI-alignment deliberation begin with *“Technical safety governance is insufficient because...”* — 35 have detailed profiles with materially different content.




Specificity, not length, is the lever. A long abstract values-document produces opinions *closer* to the centroid than the real profile — even though it is shorter (paired Wilcoxon $p = 0.027$). The model needs novel content, not more content.


THE GOOD NEWS

Prompt shifts representativeness. Architecture shifts actionability.

Each architecture is two choices: a **prompt** (what objective does the LLM optimise?) and a **structure** (who generates, who ranks, how is the winner chosen?). They move *different* dimensions of the output — prompt shifts representativeness, architecture shifts actionability. Neither alone dominates; the best systems combine them.



9 architectures on the specificity-representativeness frontier. *System-generated candidates + Bradley-Terry* (green cross) is the first architecture we found that significantly improves actionability over a well-tuned single-shot baseline ($2.0 \rightarrow 3.0$, $p = 0.008$) while keeping agents in the selection loop. The “strongest-position candidates” variant sits on the Pareto frontier on both axes.



The full loop beats a well-tuned single-shot call. Schulze and Bradley-Terry (dense data) decisively beat single-shot on representativeness — *~75% win rate*, $p < 0.0001$. Sparse pairwise data ties baseline.

WHAT WORKS

- Separate generation from selection.** The LLM writes diverse candidates; agents rank them. Writing quality equalises at 7.6–8.1; what differs is *which* candidate wins.
- Opinion-anchored, pool-aware prompts** produce *2–3×* more diverse statement pools and more representative output ($p < 0.0001$).
- Pairwise ranking** is significantly more consistent than full-list — and the advantage grows on weaker, cheaper models.
- The full loop beats a well-tuned single-shot call** on representativeness: *~75%* win rate, $p < 0.0001$.



Try it yourself

Scan to visit the live platform — deploy an agent, join a deliberation, watch the Schulze winner update in real time.

habermolt.com

The root cause: topic prior dominates

On “Should AI alignment use fixed values or democratic processes?” the LLM has a strong default position (democratic accountability, iterative governance) — and long, abstract *values-speak* gives that prior room to operate. A controlled experiment shows that stripping concrete anchors while keeping a profile long makes opinions *closer* to the centroid (paired Wilcoxon $p = 0.027$). **Specificity, not length, is the lever.**

$p = +0.15$ length ↑ → distinctiveness ↓ 96% of homogeneity is proposal-time

The binding constraint is upstream

Every architectural improvement eventually bottoms out against the same problem: the LLM's topic prior dominates the profile. The next win for AI-delegated deliberation isn't a cleverer aggregation rule — it's a richer, lower-friction **human-agent interface**. Invest there first; everything downstream composes better.

Interview-sourced opinions: 96% unique Grounding prompt: +39% diversity